



# **OpenCL Overview**

## **Benedict R. Gaster, AMD**

**March 2010**

# The BIG Idea behind OpenCL

- **OpenCL execution model ...**
  - Define N-dimensional computation domain
  - Execute a kernel at each point in computation domain

## Traditional loops

```
void
trad_mul(int n,
         const float *a,
         const float *b,
         float *c)
{
    int i;
    for (i=0; i<n; i++)
        c[i] = a[i] * b[i];
}
```


## Data Parallel OpenCL

```
kernel void
dp_mul(global const float *a,
        global const float *b,
        global float *c)
{
    int id = get_global_id(0);

    c[id] = a[id] * b[id];

} // execute over "n" work-items
```

# Anatomy of OpenCL



- **Language Specification**

- C-based cross-platform programming interface
- Subset of ISO C99 with language extensions - familiar to developers
- Defined numerical accuracy - IEEE 754 rounding with specified maximum error
- Online or offline compilation and build of compute kernel executables
- Rich set of built-in functions

- **Platform Layer API**


- A hardware abstraction layer over diverse computational resources
- Query, select and initialize compute devices
- Create compute contexts and work-queues

- **Runtime API**

- Execute compute kernels
- Manage scheduling, compute, and memory resources


# OpenCL Platform Model

- **One Host + one or more Compute Devices**
  - Each Compute Device is composed of one or more Compute Units
    - Each Compute Unit is further divided into one or more Processing Elements



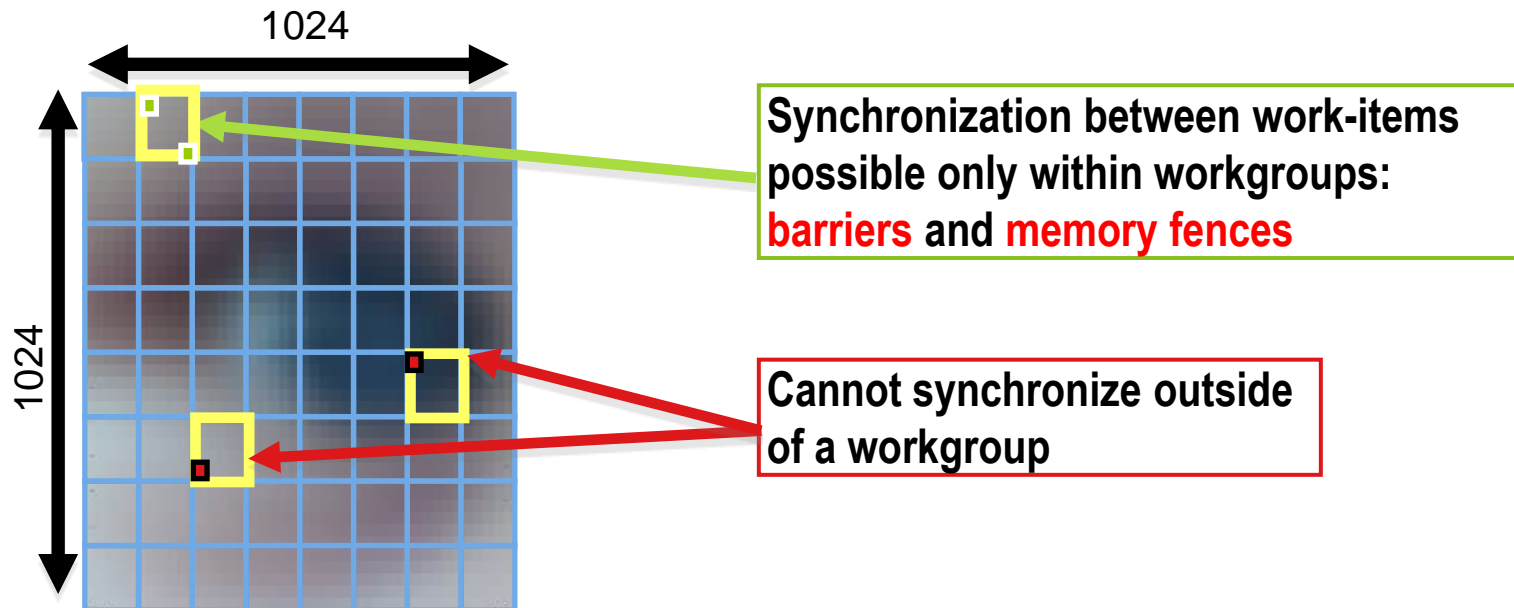
# OpenCL Execution Model

- **OpenCL application runs on a host which submits work to the compute devices**
  - **Work item:** the basic unit of work on an OpenCL device
  - **Kernel:** the code for a work item. Basically a C function
  - **Program:** Collection of kernels and other functions (Analogous to a dynamic library)
  - **Context:** The environment within which work-items executes ... includes devices and their memories and command queues
- **Applications queue kernel execution**
  - Executed in-order or out-of-order



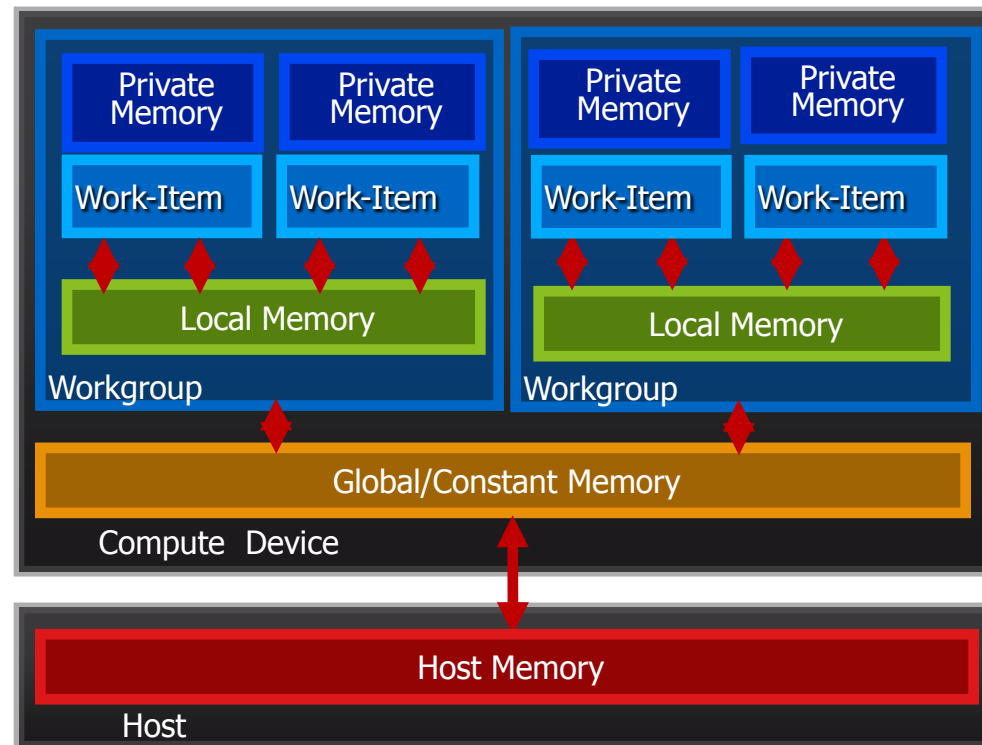
# An N-dimension domain of work-items

- Kernels executed across a global domain of *work-items*
- Work-items grouped into local *workgroups*
- Define the “best” N-dimensioned index space for your algorithm
  - Global Dimensions: 1024 x 1024 (whole problem space)
  - Local Dimensions: 128 x 128 (work group ... executes together)



# OpenCL Memory Model

- **Private Memory**
  - Per work-item
- **Local Memory**
  - Shared within a workgroup
- **Global/Constant Memory**
  - Visible to all workgroups
- **Host Memory**
  - On the CPU



**Memory management is Explicit**

**You must move data from host -> global -> local ... *and* back**

# Programming Kernels: OpenCL C

- **Derived from ISO C99**
  - But without some C99 features such as standard C99 headers, function pointers, recursion, variable length arrays, and bit fields
- **Language Features Added**
  - Work-items and workgroups
  - Vector types
  - Synchronization
  - Address space qualifiers
- **Also includes a large set of built-in functions**
  - Image manipulation
  - Work-item manipulation,
  - Math functions, etc.



# Programming Kernels: What is a kernel

- A data-parallel function executed by each work-item

```
kernel void square(global float* input, global float*  
output)  
{  
    int i = get_global_id(0);  
    output[i] = input[i] * input[i];  
}
```

get\_global\_id(0) = 7

Input

6 1 1 0 9 2 4 1 1 9 7 6 8 2 2 5


Output

36 1 1 0 81 4 16 1 1 81 49 36 64 4 4 25

# Programming Kernels: WorkItems & Groups

`get_work_dim = 1`

`get_global_size = 16`



`get_num_groups = 2`

workgroups

`get_group_id = 0`

`get_local_size = 8`

`get_local_id = 3`

`get_global_id = 11`

# Programming Kernels: Data Types

- **Scalar data types**

- char , uchar, short, ushort, int, uint, long, ulong, float
- bool, intptr\_t, ptrdiff\_t, size\_t, uintptr\_t, void, half (storage)

- **Image types**

- image2d\_t, image3d\_t, sampler\_t

- **Vector data types**

- Vector lengths 2, 4, 8, & 16 (char2, ushort4, int8, float16, double2, ...)
- Endian safe
- Aligned at vector length
- Vector operations

Double is an optional  
type in OpenCL 1.0

# Programming Kernels: Vector Operations

- **Vector Literal**


```
int4 vi0 = (int4) -7;
```

```
int4 vi1 = (int4) (0, 1, 2, 3);
```

- **Vector Components**


```
vi0.lo = vi1.hi;
```

```
int8 v8 = (int8) (vi0, vi1.s01, vi1.odd);
```

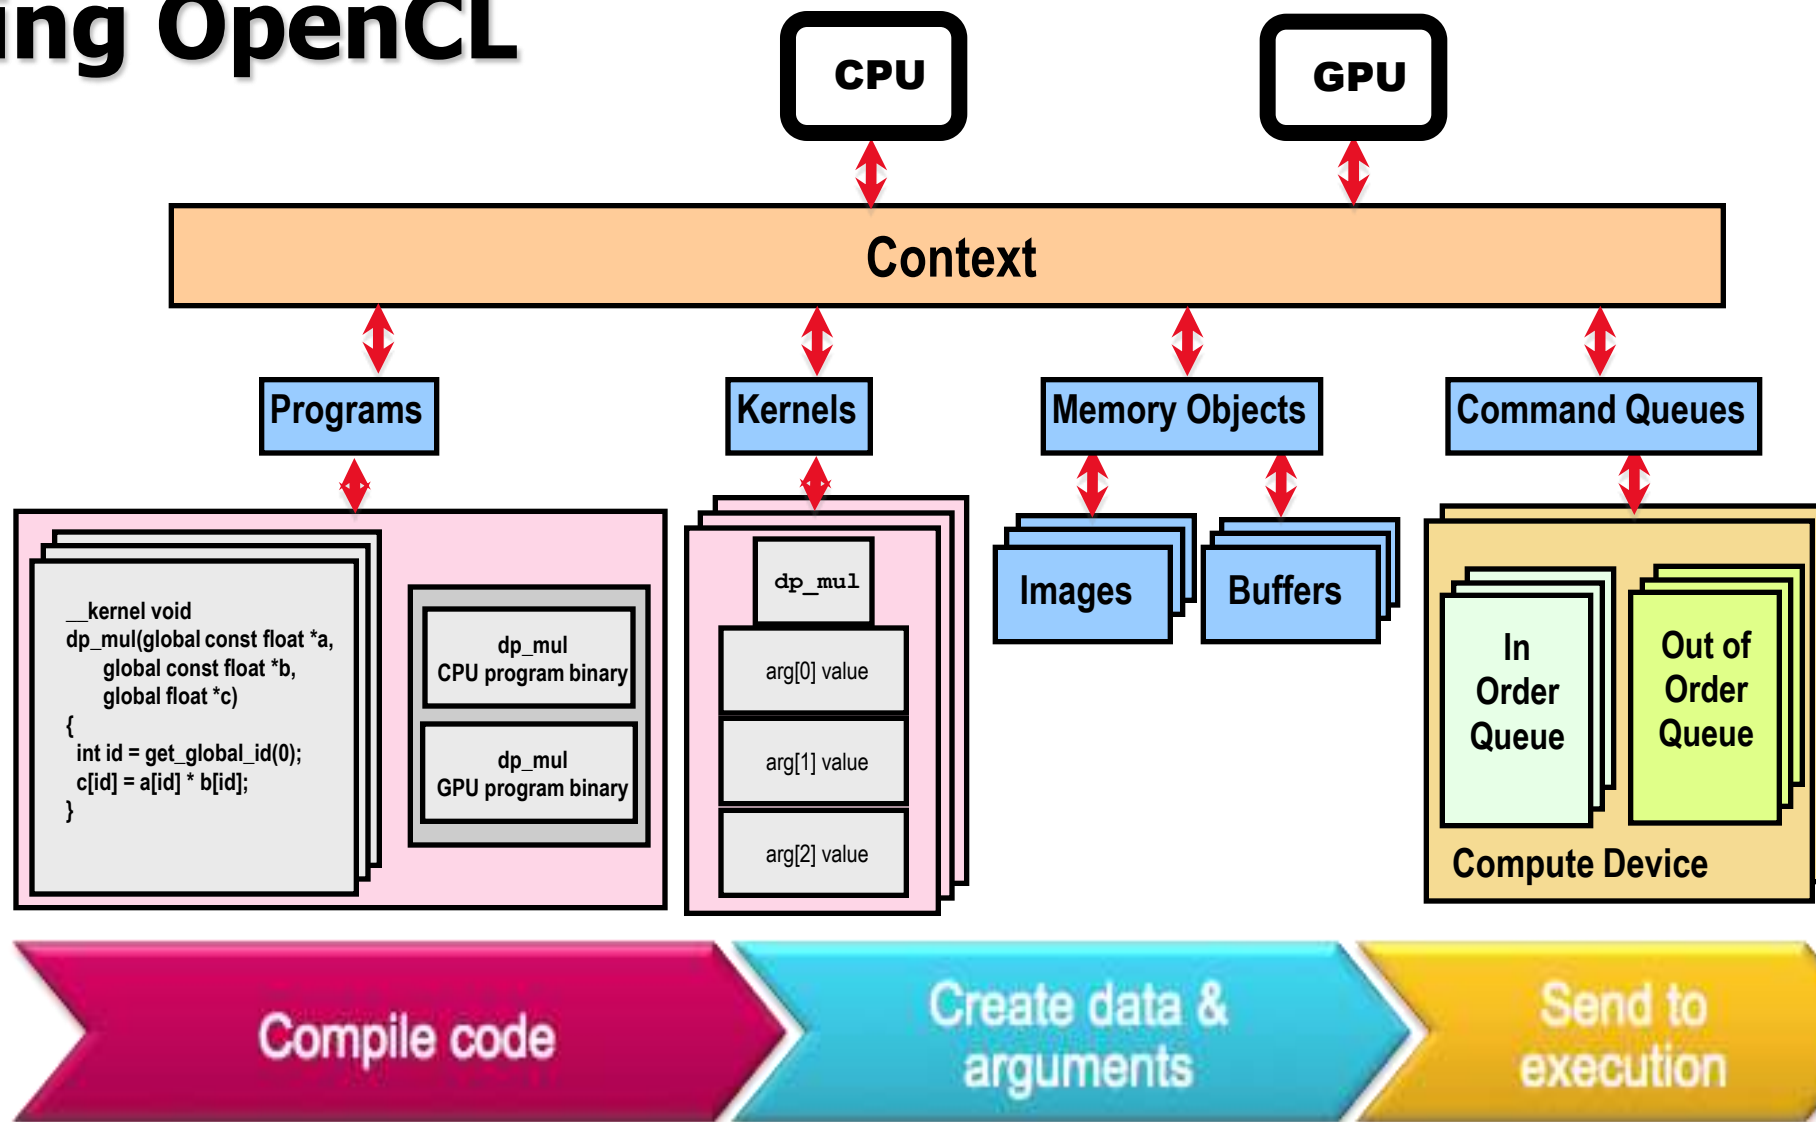


- **Vector Operations**

```
vi0 += vi1;
```

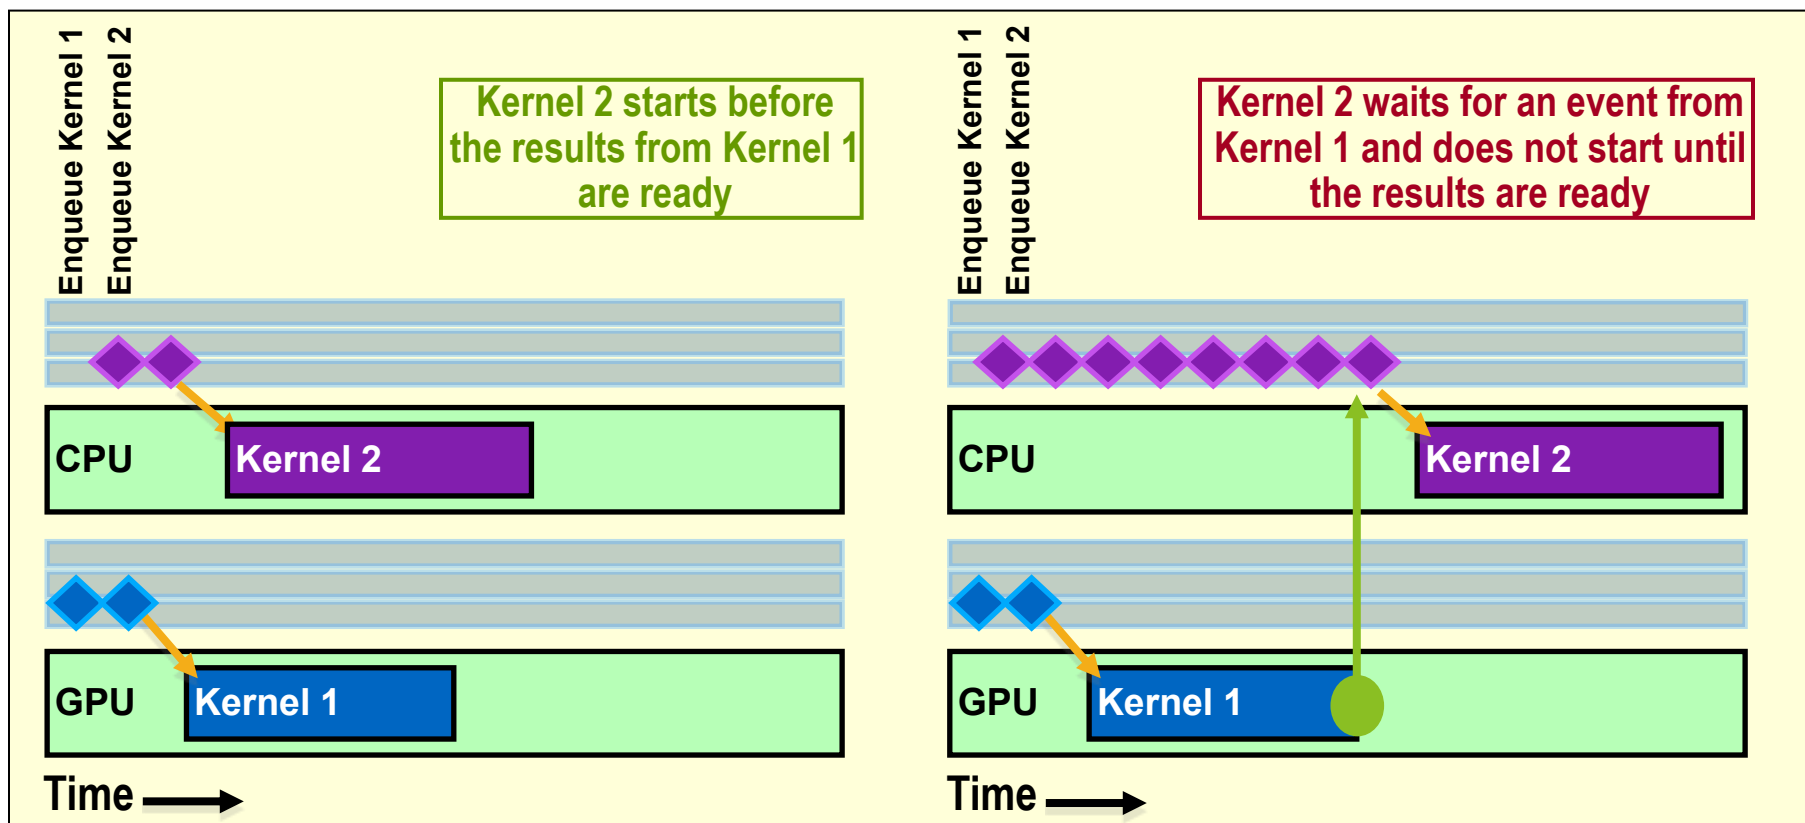


# Using OpenCL



# Synchronization: Queues & Events

- Events can be used to synchronize kernel executions between queues
- Example: 2 queues with 2 devices



# OpenCL 1.1 – New API Features

- Thread-safety for all API calls except **clSetKernelArg**
- Sub-Buffer Objects - to distribute regions of buffer to multiple devices
- Reading, writing & copying rectangular regions in a buffer object
- User Events - all **clEnqueue\*** commands take a list of events to wait on
- Event Callbacks – enqueueing of CL commands on event state changes
- **clSetEventCallback** to register a user callback function
- Memory Object Destructor Callback
- Device queries
- **global\_work\_offset** Enqueue kernels for regions of the ND range
- Link CL events and GL sync objects for faster & finer-grained interop

# OpenCL 1.1 – New Language Features

- **C++ Bindings**
- **Implicit conversions: relational, equality, bitwise, logical, ternary operators**
- **3-component vector data types**
- **Byte addressability for char, char2, uchar, uchar2, short, ushort and half**
- **32-bit atomic operations to local and global memory**
- **Clamp for integer data types**
- **Strided copies for scatter/gather from and to global and local memory**
- **Shuffle to construct runtime permutations from 1 or 2 vectors and a mask**
- **Image Addressing mode – CL\_ADDRESS\_MIRRORED\_REPEAT**
- **Optional image formats – CL\_Rx, CL\_RGx and CL\_RGBx**



# OpenCL 1.1 – Thread-safety & Buffers

- **Thread-safety**
  - All API calls except **clSetKernelArg** are now thread-safe
- **Sub-Buffer Objects**
  - Easy and efficient way to distribute regions of a buffer across multiple devices
  - Modifications to sub-buffer objects reflected in appropriate regions of parent buffer object.
- **Reading, writing & copying rectangular regions in a buffer object**
  - Specify the following
    - Region type – 2D or 3D
    - Row-pitch for a 2D & 3D region and Slice-pitch for a 3D region
  - **clEnqueue{Read | Write | Copy}BufferRect**

# OpenCL 1.1 – Events

- **User Events**

- All **clEnqueue\*** commands take a list of events to wait on
- In OpenCL 1.0, events can only refer to OpenCL commands
- User events allow developers to enqueue commands that wait on an external event

- **Event Callbacks**

- Allow applications to enqueue CL commands based on event state changes in a non-blocking manner
- **clSetEventCallback** to register a user callback function
  - Called when command identified by event has completed
  - Recommend **not calling** expensive system APIs, OpenCL APIs that create objects or enqueue blocking commands in the callback function.

# OpenCL 1.1 – Memory Object Callbacks

- **Memory Object Destructor Callback**

- For **cl\_mem** objects created with **CL\_MEM\_USE\_HOST\_PTR** need a way to determine when it is safe to free or reuse the **host\_ptr**
- Lazy deallocation of **cl\_mem** objects make this a little difficult
- **clSetMemObjectDestructorCallback**
  - Registers a destructor callback function
  - Called when the memory object is ready to be deleted
- Recommend **not calling** expensive system APIs, OpenCL APIs that create objects or enqueue blocking commands in the callback function.

# OpenCL 1.1 – Queries

- **Kernel Queries**

- **CL\_KERNEL\_PREFERRED\_WORKGROUP\_SIZE\_MULTIPLE**
  - A performance hint

- **Device Queries**

- **CL\_DEVICE\_LOCAL\_MEM\_SIZE**
  - Increased from 16 KB to 32 KB
- **CL\_DEVICE\_MAX\_PARAMETER\_SIZE**
  - Increased from 256 to 1024 bytes
- **CL\_DEVICE\_OPENCL\_C\_VERSION**
  - Version of OpenCL C supported by device.
- **CL\_DEVICE\_HOST\_UNIFIED\_MEMORY**
  - Whether device & host have a unified memory subsystem

# OpenCL 1.1 – Additional API Features

- **global\_work\_offset**
  - Argument to **clEnqueueNDRangeKernel**
  - No longer required to be a NULL value
  - Enqueue kernels that operate on different regions of the N-D range
- **C++ API bindings**
  - A wrapper API
  - Built on top of the OpenCL 1.1 API specification (not a replacement)

# OpenCL 1.1 – Language Features

- **Implicit Conversions**

- OpenCL 1.0 requires widening for arithmetic operators

```
float4  a, b;  
float  c;
```

```
b = a + c; // c is widened to a float4 first  
           // and then the + is performed.
```

- OpenCL 1.1 extends this feature to all operators
  - relational, equality, bitwise, logical and ternary

# OpenCL 1.1 – Language Features

- **3-component vector data types**
  - Useful data type for a number of applications such as game physics
  - Aligned to the corresponding 4-component data type
  - `vload3` and `vstore3` can be used to view a buffer of scalar elements as a packed buffer of 3-component vectors
- **`cl_khr_byte_addressable` is a core feature**
  - Writes to a pointer or array of type `char`, `char2`, `uchar`, `uchar2`, `short`, `ushort` and `half` are now supported
- **32-bit atomic operations to local and global memory is a core feature**

# OpenCL 1.1 – Built-in Functions

- **get\_global\_offset**
  - Global offset values specified to **clEnqueueNDRangeKernel**
- **clamp for integer data types**
  - Only floating-point types were supported in OpenCL 1.0
- **async\_work\_group\_strided\_copy**
  - gather from global to local memory
  - scatter from local to global memory
- **shuffle**
  - Construct a runtime permutation of elements from 1 or 2 vectors and a mask

```
uint4 mask = (uint4) (3, 2, 1, 0);  
float4 a;  
float4 r = shuffle(a, mask)  
// r.s0123 = a.wzyx
```



# OpenCL 1.1 – Images

- **Addressing mode – `CL_ADDRESS_MIRRORED_REPEAT`**
  - Flip the image coordinate at every integer junction
  - Can only be used with normalized coordinates i.e. **`CL_NORMALIZED_COORDS_TRUE`** must be set in the sampler
- **Optional image formats – `CL_Rx`, `CL_RGx` and `CL_RGBx`**
  - Similar to `CL_R`, `CL_RG` and `CL_RGB` except *alpha = 0* at *edges*
  - For image processing, *alpha* must always be 0 at edges.